Learn Compiler Design Through xDSL: A Practical Course for AI/ML Engineers

Djordje Todorovic

August 2025

Contents

1 Preface: Who This Course Is For and Why Compilers Matter

1.1 Prerequisites and Intended Audience

This course is designed for engineers and developers who want to understand the fundamental concepts of compiler design through hands-on experience with xDSL. To get the most out of this material, you should have:

- **Python Proficiency**: Strong knowledge of Python is essential, as xDSL is a Python-native framework. You should be comfortable with classes, decorators, type hints, and the Python standard library.

- **Basic Programming Knowledge**: Understanding of fundamental programming concepts like functions, variables, control flow, data structures, and algorithms is required.

- **Machine Learning and AI Basics**: Familiarity with ML/AI concepts is crucial - you should understand tensors, neural networks, computational graphs, and optimization techniques. Experience with frameworks like PyTorch or TensorFlow will help you appreciate the compiler optimizations we'll explore.

- **Mathematical Foundations**: Basic understanding of linear algebra and discrete mathematics will be helpful, especially when working with optimization algorithms and graph transformations.

If you're an ML engineer curious about what happens "under the hood" when your models are compiled and optimized, or a systems programmer interested in building efficient code transformation tools, this course will bridge that gap using familiar Python syntax.

1.2 Compilers: The Invisible Revolution That Changed Computing

Compilers are perhaps the most transformative technology in the history of computing, yet they work so seamlessly that we rarely think about them. To understand their revolutionary impact, consider this: the Linux kernel, which powers billions of devices worldwide, was initially written in assembly language - a tedious, error-prone process where programmers had to think in terms of individual CPU instructions and memory addresses. Each line of assembly code corresponded directly to a machine instruction, making even simple tasks require hundreds of lines of code.

Then came the evolution of the C compiler. As compilers matured, they could translate high-level C code into assembly with equal or sometimes even better performance than hand-written assembly. This transformation was revolutionary - Linux could be rewritten in C, making it portable across different architectures, easier to maintain, and accessible to a broader community of developers. What once required intimate knowledge of specific CPU architectures could now be expressed in readable, maintainable code. The compiler handled the complex translation, optimization, and architecture-specific details automatically.

This same revolution continues today in machine learning. Just as C compilers freed developers from assembly, ML compilers like XLA, TVM, and MLIR free ML engineers from writing architecture-specific kernels. Your PyTorch model can run efficiently on CPUs, GPUs, TPUs, or custom accelerators, all thanks to sophisticated compiler technology working behind the scenes.

1.3 The Core Mission: Represent and Transform

At its heart, a compiler's job is elegantly simple yet profoundly powerful: to **represent** programs in structured forms and to **transform** these representations to achieve specific goals. Think of a compiler as a sophisticated translator that not only converts between languages but also understands the meaning of what it's translating deeply enough to improve it along the way.

This representation and transformation paradigm involves:

- **Representation**: Converting source code into structured data (Abstract Syntax Trees, Intermediate Representations, Control Flow Graphs) that capture the program's semantics precisely

- **Analysis**: Understanding data dependencies, control flow, memory access patterns, and optimization opportunities within these representations

- **Transformation**: Systematically modifying these representations to optimize performance, reduce resource usage, or target specific hardware architectures

- **Preservation**: Ensuring that transformations maintain the original program's correctness and semantics

1.4   The Power of Serialize, Deserialize, and Verify

A crucial but often overlooked aspect of compiler infrastructure is the ability to serialize, deserialize, and verify intermediate representations. This capability fundamentally changes how we can work with compilers and is a cornerstone of modern compiler design.

Why These Capabilities Matter:

- **Serialization** allows us to save the compiler's intermediate state to disk, enabling separate compilation, caching of optimization results, and distribution of compiled modules. You can pause compilation, save the IR, and resume later - or on a different machine entirely.

- **Deserialization** lets us load previously compiled modules, compose them together, and build large systems incrementally. This is essential for modular compilation and linking.

- **Verification** ensures that the IR is well-formed and obeys the type system and semantic rules. This catches errors early, enables safe transformations, and provides guarantees about program behavior. Without verification, compiler bugs could silently corrupt programs.

The xDSL infrastructure we're about to explore has these capabilities built into its core. Every operation, every transformation, and every intermediate state can be serialized to a human-readable textual format, loaded back, and verified for correctness. This isn't just a convenience feature - it's fundamental to building robust, composable compiler pipelines. You can inspect the IR at any stage, debug transformations by examining the before-and-after states, and even hand-write IR for testing.

This approach contrasts sharply with traditional compilers where the intermediate states are often opaque binary structures locked inside the compiler's memory. With xDSL (inspired by MLIR), the entire compilation pipeline becomes transparent, debuggable, and extensible. You can serialize the IR after each optimization pass, analyze what changed, and even replay specific transformations. This visibility and control make compiler development more like software engineering and less like black magic.

As you progress through this course, you'll come to appreciate how these fundamental capabilities - serialize, deserialize, and verify - enable you to build reliable, maintainable, and powerful compilation tools. They transform the compiler from a monolithic black box into a modular, inspectable, and trustworthy system.

2 Introduction: Compilers as Representation and Transformation Engines

2.1 Welcome to the World of Compilers

If you're coming from an AI/ML background, you already understand the power of transforming data through layers of abstraction. Compilers are remarkably similar: they take programs written in one representation and systematically transform them into another, more optimized or more executable form.

Think of a compiler as a sophisticated pipeline that:

1. **Represents** programs as structured data (Abstract Syntax Trees, Intermediate Representations)

2. **Transforms** these representations through a series of optimization passes

3. **Analyzes** code to understand dependencies, patterns, and optimization opportunities

4. **Generates** efficient target code for specific hardware

2.2 Why xDSL?

xDSL is a Python-native compiler framework that makes compiler concepts accessible to Python programmers. Unlike traditional compiler frameworks written in C++ (like LLVM), xDSL allows you to:

- Build compilers using familiar Python syntax

- Prototype and experiment quickly

- Understand compiler internals without low-level complexity

- Leverage the entire Python ecosystem for analysis and visualization

2.3 Course Philosophy: Learn by Building

This course takes a hands-on approach. Instead of starting with theory, we'll build small compilers and optimization passes from day one. Each concept will be introduced through practical examples that you can run, modify, and experiment with.

2.4 What You'll Learn

By the end of this course, you'll understand:

- How compilers represent programs internally (IR - Intermediate Representations)

- What operations and operators mean in a compiler context

- How to write analysis passes that extract information from code

- How to implement transformations that optimize programs

- The connection between compiler optimizations and ML model optimization

3 Chapter 1: Understanding Intermediate Representations (IR)

3.1 What is an IR?

An Intermediate Representation (IR) is the compiler's internal language for representing programs. Just as neural networks use tensors to represent data, compilers use IRs to represent code structure and semantics.

Key properties of IRs:

- **Abstract**: Higher-level than machine code, but more structured than source code

- **Explicit**: Makes implicit operations visible (e.g., memory allocations, type conversions)

- **Analyzable**: Designed for easy pattern matching and transformation

- **Hierarchical**: Can represent nested structures (functions, loops, blocks)

3.2 The Compiler Pipeline: From Source to Machine Code

Before diving into xDSL specifics, let's understand how compilers work. A compiler is traditionally divided into three main phases, each with specific responsibilities:

3.2.1 Frontend: Understanding Your Code

The frontend translates source code into an intermediate representation. It consists of:
1. Lexical Analysis (Lexer/Tokenizer)

- Breaks source code into tokens (keywords, identifiers, operators)

- Like splitting a sentence into words

- Example: "x = 42 + y" becomes tokens: [ID:x, ASSIGN, NUM:42, PLUS, ID:y]

2. Syntax Analysis (Parser)

- Builds an Abstract Syntax Tree (AST) from tokens

- Checks grammatical structure

- Like diagramming a sentence to understand its structure

```
# Source: x = 42 + y
# AST representation:
Assignment(
    left=Variable("x"),
    right=BinaryOp(
        op="+",
        left=Number(42),
        right=Variable("y")
    )
)
```
Listing 1: Simple AST Example

3. Semantic Analysis

- Type checking: Is this operation valid?

- Symbol resolution: What does this variable refer to?

- Example: Can't add a string to a number, variables must be declared

3.2.2 Middle-end: Optimization Central

The middle-end works on the IR to optimize code:

- **Platform-independent**: Optimizations work for any target

- **Examples**: Dead code elimination, constant folding, loop optimization

- **Multiple passes**: Each optimization is a separate pass over the IR

3.2.3 Backend: Generating Machine Code

The backend translates optimized IR to target-specific code:

- **Instruction selection**: Choose machine instructions

- **Register allocation**: Assign variables to CPU registers

- **Instruction scheduling**: Order operations for performance

3.3 Learning from Giants: LLVM and MLIR

3.3.1 LLVM: The Industry Standard

LLVM (Low Level Virtual Machine) is the most widely used compiler infrastructure:

- Powers Clang (C/C++), Swift, Rust, and many other languages

- Provides a well-defined IR (LLVM IR) that many frontends target

- Excellent tutorial: Kaleidoscope Tutorial

- Shows how to build a complete compiler for a simple language

```
1  ; LLVM IR for: int add(int a, int b) { return a + b; }
2  define i32 @add(i32 %a, i32 %b) {
3  entry:
4    %sum = add i32 %a, %b
5    ret i32 %sum
6  }
```

Listing 2: LLVM IR Example

3.3.2 MLIR: Multi-Level IR

MLIR (Multi-Level Intermediate Representation) extends LLVM's concepts:

- Developed by Google, now part of LLVM project

- Supports multiple IRs (dialects) in the same infrastructure

- Used by TensorFlow, PyTorch, and other ML frameworks

- Tutorial: Toy Language Tutorial

- xDSL is heavily inspired by MLIR but implemented in Python

```
1  // High-level ML operation
2  %result = "tf.MatMul"(%a, %b) : (tensor<2x3xf32>, tensor<3x4xf32>)
3                                 -> tensor<2x4xf32>
4
5  // After lowering to linalg dialect
6  %result = linalg.matmul ins(%a, %b : memref<2x3xf32>, memref<3x4xf32>)
7                          outs(%c : memref<2x4xf32>)
8
9  // After lowering to loops
10 scf.for %i = 0 to 2 {
11   scf.for %j = 0 to 4 {
12     scf.for %k = 0 to 3 {
13       // actual computation
14     }
15   }
16 }
```

Listing 3: MLIR Example - Multiple Abstraction Levels

3.4 Why xDSL Uses IR Instead of AST

While traditional compilers start with AST, xDSL (like MLIR) works directly with IR because:

- **Uniformity**: All transformations work on the same structure

- **Composability**: Easy to combine different languages/dialects

- **Reusability**: Optimizations work across different source languages

- **Analysis-friendly**: SSA form makes many analyses trivial

Think of it this way:

- **AST**: Like a sentence diagram - shows structure

- **IR**: Like assembly with types - shows computation

3.5 SSA Form: Single Static Assignment

xDSL uses SSA (Single Static Assignment) form, where each variable is assigned exactly once. This is similar to functional programming and makes many analyses simpler.

```
1  # Traditional form
2  x = 5
3  x = x + 1
4  y = x * 2
5
6  # SSA form (conceptual)
7  x_0 = 5
8  x_1 = x_0 + 1
9  y_0 = x_1 * 2
```

Listing 4: Traditional vs SSA Form

Why SSA matters:

- **No ambiguity**: Each value has exactly one definition

- **Easy analysis**: Use-def chains are explicit

- **Better optimization**: Many optimizations become simpler

- **Parallel-friendly**: Dependencies are clear
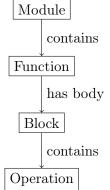
3.6 Your First xDSL Program

Let's create a simple IR program using xDSL:

```python
from xdsl.context import Context
from xdsl.dialects import builtin, arith, func
from xdsl.ir import Operation, Block, Region
from xdsl.builder import Builder

# Create a context (like a namespace for dialects)
ctx = Context()
ctx.load_dialect(builtin.Builtin)
ctx.load_dialect(arith.Arith)
ctx.load_dialect(func.Func)

# Build a simple function that adds two numbers
@Builder.implicit_region
def create_add_function():
    # Create function signature: (i32, i32) -> i32
    func_type = func.FunctionType.from_lists([builtin.i32, builtin.i32
    ], [builtin.i32])

    with func.FuncOp("add", func_type, visibility="public").body:
        # Get function arguments
        arg0 = Block.current().args[0]
        arg1 = Block.current().args[1]

        # Create add operation
        result = arith.Addi(arg0, arg1).result

        # Return the result
        func.Return(result)

module = builtin.ModuleOp([create_add_function()])
print(module)
```

Listing 5: Creating IR in xDSL

3.7 Understanding IR Structure

The IR in xDSL follows a hierarchical structure:

Module

contains

Function

has body

Block

contains

Operation

3.8 Dialects: Modular IR Design

xDSL organizes operations into *dialects* - collections of related operations and types. This is similar to how PyTorch has different modules (nn, optim, etc.).

Common dialects in xDSL:

- `arith`: Arithmetic operations (add, multiply, divide)

- `func`: Function definitions and calls

- `memref`: Memory operations (like NumPy arrays)

- `scf`: Structured control flow (loops, conditionals)

- `linalg`: Linear algebra operations (matrix multiply, convolution)

4 Chapter 2: Operations and Operators

4.1 What is an Operation?

An operation in xDSL is the fundamental unit of computation. Each operation:

- Has a name (e.g., `arith.addi` for integer addition)

- Takes input values (operands)

- Produces output values (results)

- May have attributes (compile-time constants)

- Can contain regions (nested blocks of operations)

4.2 Creating Custom Operations

Let's define a custom operation for matrix operations:

```python
from xdsl.irdl import (
    IRDLOperation,
    OpResult,
    Operand,
    OperandDef,
    ResultDef,
    op_def,
    result_def
)
from xdsl.dialects.builtin import TensorType

@op_def
class MatMulOp(IRDLOperation):
    """Matrix multiplication operation"""
    name = "my_dialect.matmul"

    # Define operands (inputs)
    lhs = operand_def(TensorType)
    rhs = operand_def(TensorType)

    # Define results (outputs)
    result = result_def(TensorType)

    def __init__(self, lhs: Value, rhs: Value):
        super().__init__(operands=[lhs, rhs],
                         result_types=[infer_matmul_type(lhs, rhs)])
```

```
28    def verify(self):
29        """Verify that the operation is well-formed"""
30        # Check that dimensions are compatible
31        lhs_shape = self.lhs.type.shape
32        rhs_shape = self.rhs.type.shape
33
34        if lhs_shape[-1] != rhs_shape[0]:
35            raise ValueError("Incompatible matrix dimensions")
```
Listing 6: Defining Custom Operations

4.3 Operation Semantics

Every operation has well-defined semantics - what it means to execute that operation:

```
1  # Arithmetic operations
2  add_op = arith.Addi(a, b)   # result = a + b
3  mul_op = arith.Muli(a, b)   # result = a * b
4
5  # Memory operations
6  alloc_op = memref.Alloc.get(shape=[10, 20], element_type=f32)  #
       Allocate 10x20 array
7  load_op = memref.Load.get(memref_value, indices)  # Load from memory
8  store_op = memref.Store.get(value, memref_value, indices)  # Store to
       memory
9
10 # Control flow operations
11 for_op = scf.For(lower_bound, upper_bound, step) # for loop
12 if_op = scf.If(condition, has_else=True)  # if-then-else
```
Listing 7: Operation Semantics Example

4.4 Attributes vs Operands

Understanding the difference between attributes and operands is crucial:

- **Attributes**: Compile-time constants (shapes, types, flags)

- **Operands**: Runtime values (variables, intermediate results)

```
1  # Constant has an attribute (the value)
2  const = arith.Constant.from_int_and_width(42, 32)   # 42 is an attribute
3
4  # Add has operands (runtime values)
5  result = arith.Addi(x, y)   # x and y are operands
6
7  # Alloc has attributes (shape) but no operands
8  mem = memref.Alloc.get(shape=[100], element_type=f32)  # shape is
       attribute
```
Listing 8: Attributes vs Operands

5 Chapter 3: Analysis Passes

5.1 What is Analysis?

Analysis passes extract information from IR without modifying it. They answer questions like:

- Which variables are used where? (Use-Def Analysis)

- Which operations can run in parallel? (Dependency Analysis)

- What are the loop bounds? (Range Analysis)

- Which operations are dead code? (Liveness Analysis)

5.2 Writing a Simple Analysis Pass

Let's write an analysis that counts operations by type:

```python
from xdsl.passes import Pass
from xdsl.ir import Operation, OpResult
from collections import defaultdict

class OperationCounterPass(Pass):
    """Count operations by type in the IR"""

    name = "count-ops"

    def apply(self, module: builtin.ModuleOp) -> None:
        op_counts = defaultdict(int)

        # Walk through all operations in the module
        for op in module.walk():
            op_counts[op.name] += 1

        # Print analysis results
        print("Operation counts:")
        for op_name, count in sorted(op_counts.items()):
            print(f"  {op_name}: {count}")

        # Store results for use by other passes
        self.op_counts = op_counts
```

Listing 9: Operation Counter Analysis

5.3 Use-Def Chains

Use-Def chains track where values are defined and used:

```python
class UseDefAnalysis(Pass):
    """Analyze use-def relationships"""

    def apply(self, module: builtin.ModuleOp) -> None:
        # Build use-def chains
        for op in module.walk():
            for result in op.results:
                print(f"Value {result} defined by {op.name}")
                print(f"  Used by: {[use.operation.name for use in
    result.uses]}")

    def find_unused_values(self, module: builtin.ModuleOp) -> list[
    OpResult]:
        """Find values that are never used"""
        unused = []
        for op in module.walk():
            for result in op.results:
```

```
16                if not result.uses:
17                    unused.append(result)
18            return unused
```

Listing 10: Use-Def Analysis

5.4 Dominance Analysis

Dominance tells us which operations must execute before others:

```
1 from xdsl.ir import Block
2 from xdsl.irdl.dominance import DominanceInfo
3
4 class DominanceAnalysis(Pass):
5     """Analyze control flow dominance"""
6
7     def apply(self, module: builtin.ModuleOp) -> None:
8         # For each function in the module
9         for func_op in module.body.ops:
10            if isinstance(func_op, func.FuncOp):
11                dom_info = DominanceInfo(func_op.body)
12
13                # Check dominance relationships
14                for block in func_op.body.blocks:
15                    dominators = dom_info.get_dominators(block)
16                    print(f"Block {block} dominated by: {dominators}")
```

Listing 11: Dominance Analysis

5.5 Dataflow Analysis

Dataflow analysis propagates information through the program:

```
1 class ConstantPropagationAnalysis(Pass):
2     """Analyze which values are compile-time constants"""
3
4     def apply(self, module: builtin.ModuleOp) -> None:
5         known_constants = {}
6
7         # First pass: identify constants
8         for op in module.walk():
9             if isinstance(op, arith.Constant):
10                known_constants[op.result] = op.value.value
11
12        # Iteratively propagate constants
13        changed = True
14        while changed:
15            changed = False
16            for op in module.walk():
17                if isinstance(op, arith.Addi):
18                    # If both operands are known constants
19                    if (op.lhs in known_constants and
20                        op.rhs in known_constants):
21                        result_value = (known_constants[op.lhs] +
22                                        known_constants[op.rhs])
23                        if op.result not in known_constants:
24                            known_constants[op.result] = result_value
25                            changed = True
26
```

```
27    return known_constants
```
Listing 12: Constant Propagation Analysis

6 Chapter 4: Transformations and Optimizations

6.1 What is a Transformation?

Transformations modify the IR to improve performance, reduce code size, or prepare for further optimizations. Unlike analysis passes, transformations actually change the program structure.
    Key principles:

- **Correctness**: Must preserve program semantics

- **Profitability**: Should improve some metric (speed, size, power)

- **Composability**: Should work well with other transformations

6.2 Pattern-Based Rewriting

xDSL provides powerful pattern matching for transformations:

```
1  from xdsl.pattern_rewriter import (
2      PatternRewriter,
3      RewritePattern,
4      op_type_rewrite_pattern
5  )
6
7  class FoldAddZero(RewritePattern):
8      """Fold x + 0 -> x"""
9
10     @op_type_rewrite_pattern
11     def match_and_rewrite(self, op: arith.Addi,
12                           rewriter: PatternRewriter) -> None:
13         # Check if right operand is zero
14         if isinstance(op.rhs.owner, arith.Constant):
15             if op.rhs.owner.value.value == 0:
16                 # Replace all uses of the add with the left operand
17                 rewriter.replace_op(op, [op.lhs])
18                 return
19
20         # Check if left operand is zero
21         if isinstance(op.lhs.owner, arith.Constant):
22             if op.lhs.owner.value.value == 0:
23                 rewriter.replace_op(op, [op.rhs])
```
Listing 13: Simple Algebraic Optimization

6.3 Common Optimizations

6.3.1 Dead Code Elimination

Remove operations whose results are never used:

```
1  class DeadCodeElimination(Pass):
2      """Remove unused operations"""
3
4      def apply(self, module: builtin.ModuleOp) -> None:
```

```
5        ops_to_remove = []
6
7        for op in module.walk():
8            # Skip operations with side effects
9            if self.has_side_effects(op):
10               continue
11
12           # Check if any results are used
13           all_dead = all(not result.uses for result in op.results)
14
15           if all_dead:
16               ops_to_remove.append(op)
17
18        # Remove dead operations
19        for op in ops_to_remove:
20            op.erase()
21
22    def has_side_effects(self, op: Operation) -> bool:
23        """Check if operation has side effects"""
24        # Memory writes, function calls, etc.
25        return isinstance(op, (memref.Store, func.Call))
```

Listing 14: Dead Code Elimination

6.3.2 Common Subexpression Elimination

Identify and eliminate redundant computations:

```
1  class CommonSubexpressionElimination(Pass):
2      """Eliminate redundant computations"""
3
4      def apply(self, module: builtin.ModuleOp) -> None:
5          # Map from (op_type, operands) to result
6          expression_map = {}
7
8          for op in module.walk():
9              if self.is_pure(op):
10                 # Create a key for this expression
11                 key = self.get_expression_key(op)
12
13                 if key in expression_map:
14                     # Found duplicate - replace with existing
15                     existing_result = expression_map[key]
16                     op.results[0].replace_by(existing_result)
17                     op.erase()
18                 else:
19                     # First occurrence - remember it
20                     expression_map[key] = op.results[0]
21
22      def get_expression_key(self, op: Operation):
23          """Create hashable key for expression"""
24          return (op.name, tuple(op.operands))
```

Listing 15: Common Subexpression Elimination

6.3.3 Loop Optimizations

Optimize loop structures for better performance:

```python
class LoopUnrolling(Pass):
    """Unroll small loops"""

    def apply(self, module: builtin.ModuleOp) -> None:
        for op in module.walk():
            if isinstance(op, scf.For):
                if self.should_unroll(op):
                    self.unroll_loop(op)

    def should_unroll(self, loop: scf.For) -> bool:
        """Decide if loop should be unrolled"""
        # Get loop bounds if constant
        if (isinstance(loop.lb.owner, arith.Constant) and
            isinstance(loop.ub.owner, arith.Constant)):
            lb = loop.lb.owner.value.value
            ub = loop.ub.owner.value.value
            iterations = ub - lb

            # Unroll small loops
            return iterations <= 4
        return False

    def unroll_loop(self, loop: scf.For):
        """Completely unroll a loop"""
        # Clone loop body for each iteration
        # Update induction variable references
        # Remove original loop
        pass  # Implementation details omitted
```

Listing 16: Loop Unrolling

6.4 Lowering Transformations

Lowering transforms high-level operations into simpler ones:

```python
class LowerMatrixOps(Pass):
    """Lower matrix operations to loops"""

    def apply(self, module: builtin.ModuleOp) -> None:
        for op in module.walk():
            if isinstance(op, MatMulOp):
                self.lower_matmul(op)

    def lower_matmul(self, matmul: MatMulOp):
        """Lower matmul to three nested loops"""
        builder = Builder.before(matmul)

        # Get dimensions
        m, k = matmul.lhs.type.shape
        _, n = matmul.rhs.type.shape

        # Allocate result
        result = builder.insert(memref.Alloc.get([m, n], f32))

        # Generate three nested loops
        with builder.insert(scf.For(0, m, 1)) as i:
            with builder.insert(scf.For(0, n, 1)) as j:
                # Initialize accumulator
```

```
24                    zero = builder.insert(arith.Constant.from_float(0.0,
       f32))
25
26                    with builder.insert(scf.For(0, k, 1, [zero])) as (k_idx
       , acc):
27                        # Load elements
28                        a_elem = builder.insert(memref.Load(matmul.lhs, [i,
        k_idx]))
29                        b_elem = builder.insert(memref.Load(matmul.rhs, [
       k_idx, j]))
30
31                        # Multiply and accumulate
32                        prod = builder.insert(arith.Mulf(a_elem, b_elem))
33                        new_acc = builder.insert(arith.Addf(acc, prod))
34
35                        scf.Yield(new_acc)
36
37                    # Store result
38                    builder.insert(memref.Store(acc, result, [i, j]))
39
40          # Replace matmul with lowered result
41          matmul.results[0].replace_by(result)
42          matmul.erase()
```

Listing 17: Lowering High-Level Operations

7 Chapter 5: Practical Examples

7.1 Building a Simple Expression Compiler

Let's build a complete compiler for arithmetic expressions:

```python
from dataclasses import dataclass
from typing import Union

# Define AST nodes
@dataclass
class Expr:
    pass

@dataclass
class Number(Expr):
    value: int

@dataclass
class BinaryOp(Expr):
    left: Expr
    op: str
    right: Expr

# Parser (simplified)
def parse(text: str) -> Expr:
    """Parse expression like '2 + 3 * 4'"""
    # Implementation omitted for brevity
    pass

# IR Generator
class ExprIRGen:
```

```python
27     def __init__(self):
28         self.builder = Builder()
29
30     def compile(self, expr: Expr) -> OpResult:
31         """Compile expression to xDSL IR"""
32         if isinstance(expr, Number):
33             return arith.Constant.from_int_and_width(
34                 expr.value, 32).result
35
36         elif isinstance(expr, BinaryOp):
37             left = self.compile(expr.left)
38             right = self.compile(expr.right)
39
40             if expr.op == '+':
41                 return arith.Addi(left, right).result
42             elif expr.op == '*':
43                 return arith.Muli(left, right).result
44             # ... other operations
45
46  # Optimizer
47  def optimize(module: builtin.ModuleOp):
48      """Apply optimization passes"""
49      passes = [
50          ConstantFolding(),
51          CommonSubexpressionElimination(),
52          DeadCodeElimination(),
53      ]
54
55      for pass_instance in passes:
56          pass_instance.apply(module)
57
58  # Complete pipeline
59  def compile_expression(text: str):
60      # Parse
61      ast = parse(text)
62
63      # Generate IR
64      ir_gen = ExprIRGen()
65      result = ir_gen.compile(ast)
66
67      # Wrap in module
68      module = builtin.ModuleOp([
69          func.FuncOp("compute",
70                      func.FunctionType.from_lists([], [builtin.i32]),
71                      body=[func.Return(result)])
72      ])
73
74      # Optimize
75      optimize(module)
76
77      return module
```

Listing 18: Expression Compiler

7.2 Implementing a Domain-Specific Optimization

Let's create an optimization for neural network operations:

```
1  class FuseBatchNormIntoConv(Pass):
2      """Fuse batch normalization into preceding convolution"""
3
4      def apply(self, module: builtin.ModuleOp) -> None:
5          for conv in module.walk():
6              if isinstance(conv, nn.Conv2d):
7                  # Look for batch norm following conv
8                  for use in conv.result.uses:
9                      if isinstance(use.operation, nn.BatchNorm2d):
10                         self.fuse(conv, use.operation)
11
12     def fuse(self, conv: nn.Conv2d, bn: nn.BatchNorm2d):
13         """Fuse batch norm parameters into convolution"""
14         # Extract batch norm parameters
15         gamma = bn.gamma  # scale
16         beta = bn.beta    # shift
17         mean = bn.running_mean
18         var = bn.running_var
19         eps = bn.eps
20
21         # Compute fused weights and bias
22         std = sqrt(var + eps)
23         fused_weight = conv.weight * (gamma / std)
24         fused_bias = gamma * (conv.bias - mean) / std + beta
25
26         # Create new convolution with fused parameters
27         fused_conv = nn.Conv2d(
28             weight=fused_weight,
29             bias=fused_bias,
30             # ... other conv parameters
31         )
32
33         # Replace conv -> bn with fused conv
34         bn.result.replace_by(fused_conv.result)
35         conv.erase()
36         bn.erase()
```
Listing 19: Neural Network Optimization

7.3 Building a Mini ML Compiler

Let's create a compiler for a subset of ML operations:

```
1  # Define ML dialect
2  @dialect_def
3  class MLDialect(Dialect):
4      name = "ml"
5
6  @op_def
7  class LinearOp(IRDLOperation):
8      """Linear layer: y = Wx + b"""
9      name = "ml.linear"
10
11     input = operand_def(TensorType)
12     weight = operand_def(TensorType)
13     bias = operand_def(TensorType)
14
15     output = result_def(TensorType)
```

```
16
17 @op_def
18 class ReluOp(IRDLOperation):
19     """ReLU activation"""
20     name = "ml.relu"
21
22     input = operand_def(TensorType)
23     output = result_def(TensorType)
24
25 # Optimization: Fuse Linear + ReLU
26 class FuseLinearRelu(RewritePattern):
27     @op_type_rewrite_pattern
28     def match_and_rewrite(self, relu: ReluOp,
29                           rewriter: PatternRewriter):
30         # Check if input is from linear
31         if isinstance(relu.input.owner, LinearOp):
32             linear = relu.input.owner
33
34             # Create fused operation
35             fused = LinearReluOp(
36                 linear.input,
37                 linear.weight,
38                 linear.bias
39             )
40
41             # Replace both ops with fused version
42             rewriter.replace_op(relu, [fused.output])
43             rewriter.erase_op(linear)
44
45 # Lower to hardware operations
46 class LowerMLToHardware(Pass):
47     """Lower ML ops to target hardware"""
48
49     def apply(self, module: builtin.ModuleOp):
50         target = self.get_target()  # CPU, GPU, TPU, etc.
51
52         for op in module.walk():
53             if isinstance(op, LinearOp):
54                 if target == "gpu":
55                     self.lower_linear_to_cublas(op)
56                 elif target == "cpu":
57                     self.lower_linear_to_loops(op)
```

Listing 20: Mini ML Compiler

8 Chapter 6: Advanced Topics

8.1 Multi-Level IR and Progressive Lowering

Real compilers use multiple levels of abstraction:

```
1 class CompilationPipeline:
2     """Progressive lowering through multiple IR levels"""
3
4     def compile(self, source_code: str) -> str:
5         # Level 1: High-level ML operations
6         ml_ir = parse_to_ml_dialect(source_code)
7         ml_ir = optimize_ml_level(ml_ir)
```

```
8
9      # Level 2: Linear algebra operations
10     linalg_ir = lower_ml_to_linalg(ml_ir)
11     linalg_ir = optimize_linalg_level(linalg_ir)
12
13     # Level 3: Loops and memory operations
14     loop_ir = lower_linalg_to_loops(linalg_ir)
15     loop_ir = optimize_loops(loop_ir)
16
17     # Level 4: Target-specific operations
18     target_ir = lower_to_target(loop_ir)
19     target_ir = optimize_target_specific(target_ir)
20
21     # Code generation
22     return generate_code(target_ir)
```

Listing 21: Multi-Level Compilation Pipeline

8.2   Cost Models and Autotuning

Choosing the best optimization requires understanding costs:

```python
1  class CostModel:
2      """Estimate operation costs"""
3
4      def estimate_cost(self, op: Operation) -> float:
5          if isinstance(op, arith.Mulf):
6              return 1.0  # FP multiply cost
7          elif isinstance(op, memref.Load):
8              return self.memory_latency(op)
9          elif isinstance(op, scf.For):
10             iterations = self.analyze_trip_count(op)
11             body_cost = sum(self.estimate_cost(body_op)
12                             for body_op in op.body.walk())
13             return iterations * body_cost
14
15 class AutotuningPass(Pass):
16     """Try different optimization strategies"""
17
18     def apply(self, module: builtin.ModuleOp):
19         strategies = [
20             self.try_vectorization,
21             self.try_loop_tiling,
22             self.try_parallelization,
23         ]
24
25         best_module = module
26         best_cost = self.cost_model.estimate_cost(module)
27
28         for strategy in strategies:
29             candidate = strategy(module.clone())
30             cost = self.cost_model.estimate_cost(candidate)
31
32             if cost < best_cost:
33                 best_module = candidate
34                 best_cost = cost
35
```

```
36        return best_module
```
Listing 22: Cost-Based Optimization

8.3 Connection to ML Compilation

Modern ML frameworks are essentially compilers:

- **PyTorch JIT**: Traces Python code to TorchScript IR

- **TensorFlow XLA**: Compiles TF graphs to optimized kernels

- **JAX**: Uses XLA for JIT compilation

- **Apache TVM**: Optimizes ML models for various hardware

The concepts you've learned apply directly:

- **Graph optimization** = IR transformation

- **Kernel fusion** = Operation merging

- **Memory planning** = Register allocation

- **Quantization** = Type lowering

9 Chapter 7: Hands-On Exercises

9.1 Exercise 1: Build Your Own Optimizer

Implement an optimization that eliminates redundant memory operations:

```
1 class StoreLoadForwarding(Pass):
2     """Forward stored values to subsequent loads"""
3
4     def apply(self, module: builtin.ModuleOp):
5         # TODO: Your implementation here
6         # Hint: Track recent stores to each memory location
7         # Replace loads with stored values when possible
8         pass
```
Listing 23: Exercise: Store-Load Forwarding

9.2 Exercise 2: Create a Custom Dialect

Design a dialect for quantum computing operations:

```
1 @dialect_def
2 class QuantumDialect(Dialect):
3     name = "quantum"
4
5 # TODO: Define operations like:
6 # - Hadamard gate
7 # - CNOT gate
8 # - Measurement
9 # - Quantum circuit optimization passes
```
Listing 24: Exercise: Quantum Dialect

9.3 Exercise 3: Implement Vectorization

Transform scalar operations to vector operations:

```python
class Vectorization(Pass):
    """Automatically vectorize loops"""

    def apply(self, module: builtin.ModuleOp):
        # TODO: Find vectorizable loops
        # Transform scalar ops to vector ops
        # Handle loop remainder
        pass
```

Listing 25: Exercise: Auto-Vectorization

9.4 Exercise 4: Build a Simple JIT Compiler

Create a JIT compiler for mathematical expressions:

```python
class ExpressionJIT:
    """JIT compile and execute expressions"""

    def compile(self, expr_string: str):
        # Parse expression
        # Generate IR
        # Optimize
        # Generate machine code (or Python bytecode)
        # Return callable function
        pass

# Usage:
jit = ExpressionJIT()
f = jit.compile("x * 2 + y")
result = f(x=5, y=3)  # Should return 13
```

Listing 26: Exercise: JIT Compiler

10 Chapter 8: Real-World Applications

10.1 Case Study 1: Optimizing Neural Network Inference

Let's optimize a real neural network for deployment:

```python
def optimize_nn_for_deployment(model: NeuralNetwork):
    """Optimize neural network for inference"""

    # Convert to IR
    ir = convert_nn_to_ir(model)

    # Apply optimizations
    optimizations = [
        # Structural optimizations
        FuseLayers(),             # Conv+BN, Linear+ReLU
        RemoveDropout(),          # Not needed for inference

        # Quantization
        QuantizeWeights(bits=8),
        QuantizeActivations(bits=8),

```

```
17          # Memory optimizations
18          InplaceOperations(),     # Reuse buffers
19          MemoryPlanning(),        # Minimize allocation
20
21          # Target-specific
22          UseTargetIntrinsics(),   # AVX, NEON, etc.
23          AutoTuneKernels(),       # Platform-specific tuning
24      ]
25
26      for opt in optimizations:
27          ir = opt.apply(ir)
28
29      return ir
```

Listing 27: Neural Network Optimization Pipeline

10.2   Case Study 2: Domain-Specific Language Compiler

Build a compiler for financial computations:

```
1  # Define financial operations dialect
2  @dialect_def
3  class FinanceDialect(Dialect):
4      name = "finance"
5
6      @op_def
7      class BlackScholesOp(IRDLOperation):
8          """Black-Scholes option pricing"""
9          name = "finance.black_scholes"
10
11          spot = operand_def(f64)
12          strike = operand_def(f64)
13          time = operand_def(f64)
14          rate = operand_def(f64)
15          volatility = operand_def(f64)
16
17          call_price = result_def(f64)
18          put_price = result_def(f64)
19
20  # Optimize for numerical stability
21  class NumericalStabilityPass(Pass):
22      """Ensure numerical stability in financial calculations"""
23
24      def apply(self, module: builtin.ModuleOp):
25          for op in module.walk():
26              if isinstance(op, math.ExpOp):
27                  # Prevent overflow/underflow
28                  self.add_range_check(op)
29              elif isinstance(op, arith.DivfOp):
30                  # Prevent division by zero
31                  self.add_zero_check(op)
```

Listing 28: Financial DSL Compiler

10.3   Case Study 3: Compiler for Distributed Computing

Compile programs for distributed execution:

```python
1  @dialect_def
2  class DistributedDialect(Dialect):
3      name = "dist"
4
5      @op_def
6      class AllReduceOp(IRDLOperation):
7          """All-reduce across distributed nodes"""
8          name = "dist.allreduce"
9
10         input = operand_def(TensorType)
11         output = result_def(TensorType)
12         reduction = attr_def(str)  # "sum", "max", etc.
13
14 class DistributionPass(Pass):
15     """Distribute computation across nodes"""
16
17     def apply(self, module: builtin.ModuleOp):
18         # Analyze data dependencies
19         deps = self.analyze_dependencies(module)
20
21         # Partition computation
22         partitions = self.partition_computation(module, deps)
23
24         # Insert communication operations
25         for partition in partitions:
26             self.insert_communication_ops(partition)
27
28         # Generate distributed code
29         return self.generate_distributed_code(partitions)
```

Listing 29: Distributed Computing Compiler

11 Conclusion: Your Journey Forward

11.1 What You've Learned

Congratulations! You now understand:

- How compilers represent programs as structured data (IR)

- The role of operations, attributes, and types

- How to write analysis passes to understand code

- How to implement transformations that optimize programs

- The connection between compiler techniques and ML optimization

11.2 Next Steps

Your compiler journey doesn't end here:

1. **Explore Real Compilers**: Study LLVM, GCC, or V8

2. **Contribute to xDSL**: Add new dialects or optimizations

3. **Apply to Your Domain**: Build domain-specific compilers

4. **Learn Advanced Topics**:

- Polyhedral compilation
- Verified compilation
- Quantum compilation
- Neuromorphic compilation

11.3   Resources for Further Learning

- **xDSL Documentation**: `https://xdsl.dev`

- **MLIR Tutorial**: `https://mlir.llvm.org/docs/Tutorials/`

- **Engineering a Compiler** (Cooper & Torczon)

- **Modern Compiler Implementation** (Appel)

- **The SSA Book**: `http://ssabook.gforge.inria.fr/`

11.4   Final Thoughts

Compilers are the bridge between human intent and machine execution. As AI/ML systems become more complex and heterogeneous hardware becomes more prevalent, understanding compilation becomes crucial. The techniques you've learned - representation, analysis, and transformation - are fundamental to building efficient, scalable systems.

Remember: Every optimization you write makes programs faster, every analysis you implement makes systems more reliable, and every transformation you create pushes the boundaries of what's computationally possible.

Happy compiling!

A   Appendix A: xDSL Quick Reference

A.1   Common Operations

```
# Arithmetic
arith.Constant.from_int_and_width(42, 32)   # Integer constant
arith.Constant.from_float(3.14, f32)        # Float constant
arith.Addi(a, b)   # Integer addition
arith.Muli(a, b)   # Integer multiplication
arith.Addf(x, y)   # Float addition
arith.Mulf(x, y)   # Float multiplication

# Memory
memref.Alloc.get(shape=[10, 20], element_type=f32)
memref.Load(mem, indices)
memref.Store(value, mem, indices)

# Control Flow
scf.For(start, end, step, init_values)
scf.If(condition, has_else=True)
scf.While(init_values)

# Functions
func.FuncOp(name, func_type, body)
func.Call(func_name, args)
func.Return(values)
```

Listing 30: xDSL Operations Quick Reference

A.2 Common Patterns

```python
# Walking operations
for op in module.walk():
    process(op)

# Pattern matching
if isinstance(op, arith.Addi):
    handle_add(op)

# Building IR
with Builder.implicit_region():
    result = arith.Addi(a, b).result

# Replacing operations
old_op.results[0].replace_by(new_value)
old_op.erase()
```

Listing 31: Common xDSL Patterns

B Appendix B: Setting Up Your Environment

B.1 Installation

```bash
# Install xDSL
pip install xdsl

# For development
git clone https://github.com/xdslproject/xdsl.git
cd xdsl
pip install -e ".[dev]"

# Run tests
pytest tests/
```

Listing 32: Installing xDSL

B.2 Your First xDSL Program

```python
#!/usr/bin/env python3

from xdsl.context import Context
from xdsl.dialects import builtin, arith, func
from xdsl.printer import Printer

# Create context and load dialects
ctx = Context()
ctx.load_dialect(builtin.Builtin)
ctx.load_dialect(arith.Arith)
ctx.load_dialect(func.Func)

# Build a simple function
func_type = func.FunctionType.from_lists([builtin.i32], [builtin.i32])

with func.FuncOp("double", func_type).body as func_body:
    arg = func_body.block.args[0]
```

```
18      two = arith.Constant.from_int_and_width(2, 32)
19      result = arith.Muli(arg, two)
20      func.Return(result)
21
22  # Create module
23  module = builtin.ModuleOp([func_body.owner])
24
25  # Print the IR
26  printer = Printer()
27  printer.print_op(module)
```

Listing 33: hello*compiler.py*

C   Appendix C: Glossary

AST Abstract Syntax Tree - Tree representation of source code structure

Basic Block Sequence of operations with single entry and exit

CFG Control Flow Graph - Graph showing possible execution paths

Dataflow Analysis tracking how data moves through program

Dialect Collection of related operations and types in xDSL

Dominance Relationship where one operation must execute before another

IR Intermediate Representation - Compiler's internal program representation

Lowering Transforming high-level operations to simpler ones

Operation Fundamental unit of computation in xDSL

Operand Input value to an operation

Pass Transformation or analysis applied to IR

Pattern Template for matching and replacing IR structures

Region Container for blocks in xDSL (e.g., function body, loop body)

Result Output value from an operation

SSA Single Static Assignment - Form where each variable assigned once

Type Description of data format and constraints

Use-Def Chain Links between value definitions and uses

Value Data flowing between operations (operands and results)

Verification Checking that IR is well-formed and valid